

Multiclass classification

Multiclass strategies

- Reduce to binary

One-vs-All

$$Y = \{1, \dots, K\}$$

foreach k in Y :

$$z_i = 1 \leftrightarrow y_i == k$$

$$f_k = \text{clf.train}(X, z)$$

$$y_p(x) = \operatorname{argmax}_k f_k(x)$$

- Train a classifier to distinguish each label from the rest
- Total K classifiers
- Prediction: pick label with the highest score

One-vs-One

- Train independent classifier for each pair of labels (i, j)
- Total: $K(K-1)/2$ classifiers
- Prediction: select most common class via voting

$$y_p(x) = \operatorname{argmax}_k \sum_{i=1}^k f_{ik}(x)$$

Multiclass strategies

One-vs-All

- Linear in num. of classes

One-vs-One

- Quadratic in num. of classes

Multiclass strategies

One-vs-All

- Linear in num. of classes
- Unbalanced classes problem
- Biases in base models' scores

One-vs-One

- Quadratic in num. of classes
- Works faster with non-scalable models (e.g. kernels)
- More ambiguity in predictions

Multiclass strategies

- Reduction to binary
- Model extension

Multiclass metrics

Recap: binary classification

$$y_i \in \{0, 1\}$$

$$a_i \in \{0, 1\}$$

Error type	Prediction	Ground truth
True Positive (TP)	1	1
True Negative (TN)	0	0
False Positive (FP)	1	0
False Negative (FN)	0	1

Multiclass metrics

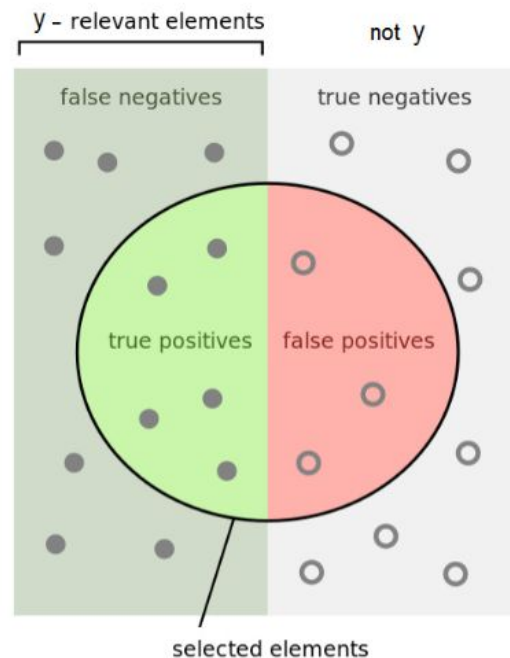
- Let's extend binary classification metrics

for each class $y \in Y$:

TP_y - True positive predictions

FP_y - False positive predictions

FN_y - False negative predictions



Micro-averaging

$$\text{Precision: } P = \frac{\sum_y \text{TP}_y}{\sum_y (\text{TP}_y + \text{FP}_y)};$$

$$\text{Recall: } R = \frac{\sum_y \text{TP}_y}{\sum_y (\text{TP}_y + \text{FN}_y)};$$

Micro-averaging

$$\text{Precision: } P = \frac{\sum_y \text{TP}_y}{\sum_y (\text{TP}_y + \text{FP}_y)};$$

$$\text{Recall: } R = \frac{\sum_y \text{TP}_y}{\sum_y (\text{TP}_y + \text{FN}_y)};$$

- Does not cover imbalanced classes

Macro-averaging

$$\text{Precision: } P = \frac{1}{|Y|} \sum_y \frac{TP_y}{TP_y + FP_y};$$

$$\text{Recall: } R = \frac{1}{|Y|} \sum_y \frac{TP_y}{TP_y + FN_y};$$

- Just average class scores

Micro-averaging of macro-averaging?

4 classes; model always outputs 1

C1: TP = 1, FP = 0

C2: TP = 1, FP = 0

C3: TP = 53, FP = 47

C4: TP = 1, FP = 1

Micro-averaging of macro-averaging?

4 classes; model always outputs 1

C1: TP = 1, FP = 0

C2: TP = 1, FP = 0

C3: TP = 53, FP = 47

C4: TP = 1, FP = 1

Precision_{micro} = 0.53

Precision_{macro} = 0.76

Feature selection

Feature selection

1. Model-free (statistical)
2. Model-based (intrinsic)
3. Performance-based

Statistical methods

- Correlation

$$R_j = \frac{\sum_{i=1}^{\ell} (x_{ij} - \bar{x}_j)(y_i - \bar{y})}{\sqrt{\sum_{i=1}^{\ell} (x_{ij} - \bar{x}_j)^2 \sum_{i=1}^{\ell} (y_i - \bar{y})^2}}$$

Statistical methods

- T-score (binary classification)

$$R_j = \frac{|\mu_0 - \mu_1|}{\sqrt{\frac{\sigma_0^2}{n_0} + \frac{\sigma_1^2}{n_1}}},$$

Statistical methods

- F-score (multiclass)

$$R_j = \frac{\sum_{k=1}^K \frac{n_j}{K-1} (\mu_j - \mu)^2}{\frac{1}{\ell-K} \sum_{k=1}^K (n_j - 1) \sigma_j^2},$$

Model-based

- Linear model

$$h(x, w) = x^T w = \sum x_i w_i$$

Model-based

- Linear model

$$h(x, w) = x^T w = \sum x_i w_i$$

- Weights are proportional to corresponding features' impact on prediction
- Do not forget about feature scaling!

Model-based

- Decision trees
- **Decrease in impurity**

$$Imp(j) = \sum_t I\{j_t = j\} p(t) \Delta_i(t)$$

$$\Delta_i(t) = i(t) - \frac{N(t_L)}{N} i(T_L) - \frac{N(t_R)}{N} i(T_R)$$

Model-based

- Random Forest
- **Mean decrease impurity**
- **Out-of-bag score**

$$\text{OOB} = \sum_{i=1}^{\ell} L\left(y_i, \frac{1}{\sum_{n=1}^N [x_i \notin X_n]} \sum_{n=1}^N [x_i \notin X_n] b_n(x_i)\right)$$

Performance-based

- Train model on various subsets and select those which perform best
- Estimate quality on hold-out set for not overfitting

J - a set of features, $\|J\| = j$

μ_J - model trained only on J parameters

$$Q(J) = Q(\mu_J, X_{test})$$

$$Q(J) \rightarrow \min$$

Performance-based

Full Search

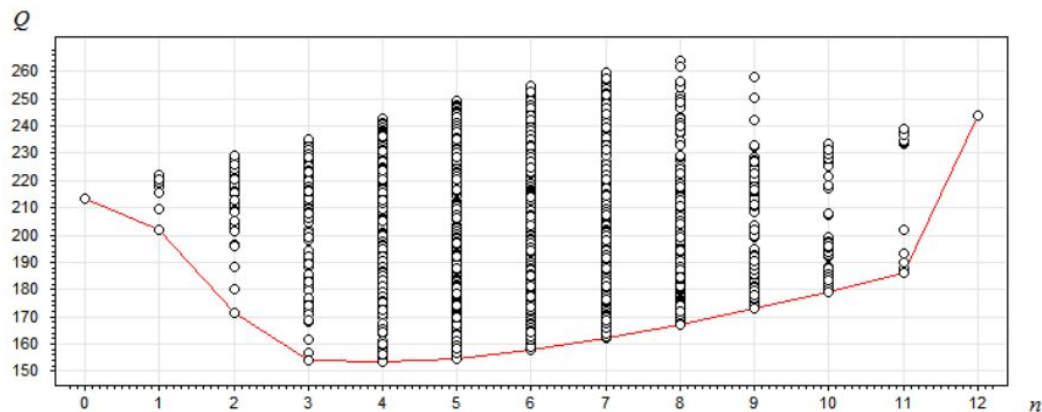
$$J = \emptyset$$

for j in $1..n$:

$$J = \operatorname{argmin}_{\|J\|=j} Q(\mu_J, X_{\text{test}})$$

if $Q(\mu_J) < Q^* : J^* = J; Q^* = Q(\mu_J)$

if $\|J\| > \|J^*\| + d : \text{return } J^*$



Performance-based

Full search

Pros:

- Simplicity
- Optimal solution

Cons:

- $O(2^n)$
- Prone to overfitting

Performance-based

Greedy addition

$$J = \emptyset$$

for j *in* $1..n$:

$$f_j = \operatorname{argmin}_{f \in F \setminus J} Q(\mu_J \cup f, X_{\text{test}})$$

$$J = J \cup f_j$$

$$\text{if } Q(\mu_J) < Q^* : J^* = J; Q^* = Q(\mu_J)$$

$$\text{if } j > \|J^*\| + d : \text{return } J^*$$

Performance-based

Greedy addition

Pros:

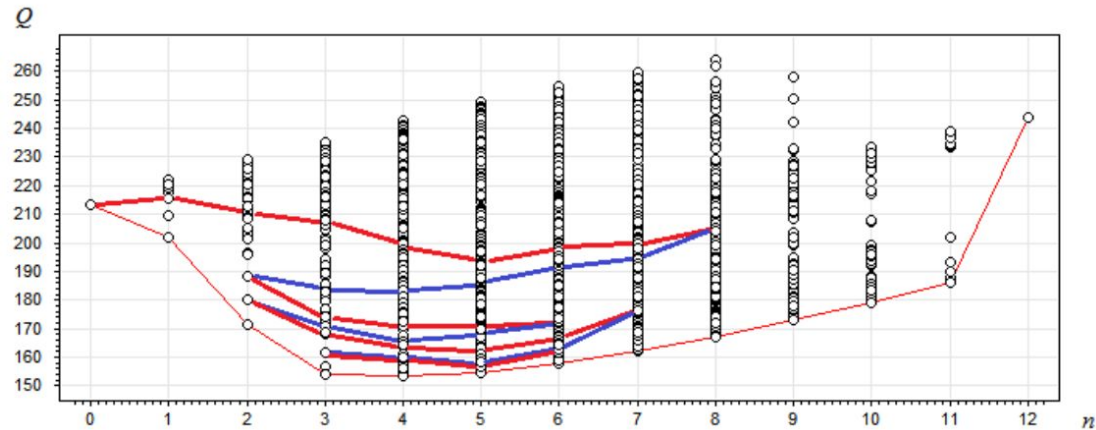
- $O(n^2)$
- Fast incremental algorithms (step-wise regression)

Cons:

- Tends to include odd features

Performance-based

Greedy addition / deletion (add-del)



Summary

- Multiclass classification
 - Classification strategies (One-vs-All, One-vs-One)
 - Metrics averaging
- Feature selection
 - Statistical
 - Model-based (intrinsic)
 - Performance-based

Summary

The following awesome materials were used:

- ML lectures by K.Vorontsov: [materials](#)
- HSE ML course by E.Sokolov: [materials](#)

Thank you for your attention!